

CYBORG NEWS

CYBORGNEWS.MIT.EDU

WEDNESDAY, SEPTEMBER 16, 2026 · VOL. II

FREE TO HUMANS

SEPT MODELS
– 12 released 8 providers

ATRIA DAWN
▲ Sept 12 release latest preview

CLAUDE FABLE 5.1
– Sept 1 ship Anthropic

MYTHOS 5.1
– Sept 1 twin trusted-access

GEMINI 3.8 FLASH
▲ Sept 2 launch Google DeepMind

GPT-6 ASTRA
▲ Sept 3 release OpenAI ship

MUSE SPARK 1.3
▼ \$0.10 per M tokens blended price

What's News

Also noted

Brain geometry shapes learning Busch and colleagues show intrinsic activity geometry keeps learning inside existing neural manifolds.

Survey warns out-sourced thinking atrophies Very Big Brain synthesis argues use dependent skills fade when AI assistants consistently bypass them.

Editorial links reasoning theory mind PMC editorial frames shared AI neuroscience problem of learning uncertainty and self to other modeling.

Fifty five cognition pre-prints listed arXiv September neurons and cognition roundup indexes 55 titles from sleep rhythms to cortical geometry without synthesis.

Review probes chatbot executive function MDPI systematic review asks whether chatbot interaction reinforces planning and goal directed control.

In brief

Gateway lists September releases Timeline logs Atria Dawn Preview September 12 plus Fugu models September 11 and DeepSeek Flash September 10.

White House oversight mechanisms expand Hinshaw notes June 2 order formalizes federal role without direct business duties amid Mythos and Fable export worries.

Dated ledger tracks breaking changes DigitalApplied dates Claude Fable 51 and Mythos 51 September 1 with three API breaks plus Gemini 38 Flash and Muse Spark 13 September 2.


Project Lily reviews private chats 404 Media reports Crossing Hurdles contractors via Mercor score real ChatGPT chats with personal details still inside.

California frontier law bites Wikipedia notes TFAIA in effect January 1 2026 plus Newsom September 10 child chatbot risk assessment laws with fines for harm.

WASHINGTON

Congress stalls as leaders reject AI slowdown plea

The House has one week before a pre-midterm break, with no safety bill close to a vote.



STALLED Visitors stand on the steps of the United States Capitol on a clear day. No AI safety measure is expected to reach a vote before the midterm break.

BALON GREYJOY · CC BY-SA 4.0

WASHINGTON — Congress is not expected to pass artificial intelligence safety legislation before the midterm elections after House Speaker Mike Johnson and President Donald Trump rejected a slowdown plea from Dario Amodei, Sam Altman, Demis Hassabis and Elon Musk. The plea followed the resignation of an Anthropic researcher who warned on X that AI could “kill us all in a decade”. The House has the week of Sept. 14 before a break until after the elections.

Johnson said Tuesday that industry concerns about AI wiping out humanity are

overstated and that companies can self-regulate. He said “But we cannot have a moratorium on the development of AI, okay, because then we will lose our edge to China.” Trump wrote Sept. 14 on Truth Social that industry leaders calling for regulation was suspect and called claims that AI will destroy the world a “hoax”. Johnson said he and Trump “are summoning” AI leaders to the White House, possibly within the next week.

“But we cannot have a moratorium on the development of AI, okay, because then we will lose our edge to China.”

MIKE JOHNSON, HOUSE

Amodei published a Sept. 12 essay urging a slower pace of developing AI capability. Altman, Hassabis and Musk agreed with him. Senate Majority Leader John Thune said he continues work with Senator Amy Klobuchar on an AI bill but added that traction before the elections is uncertain.

Senate Majority Whip John Barrasso said discussions are ongoing but that he does not see a proposal reaching a vote given the tight timeline and other legislative priorities.

In the House, Representatives Jay Obernolte and Lori Trahan introduced the Frontier Act in July to require audits and assessments. Representatives Ted Lieu and Nathaniel Moran introduced a kill switch bill. Senator Bernie Sanders and Representative Greg Casar announced plans on Sept. 3 for a permanent superintelligence ban.

GLOBAL RULES

EU Starts Audits as California Brazil India Decide

September brings statutory deadlines for high-risk systems, frontier models and national AI liability rules.

Global AI regulation shifts from preparation to enforcement in September. The European AI Office begins audits of high-risk systems. California faces a September 30 decision on frontier model safety legislation. Brazil schedules a Senate vote on comprehensive AI legislation. India introduces a liability framework for generative AI outputs.

Transition period for general high-risk systems concluded on August 2. The European AI Office in Brussels will work with 24 national market surveillance authorities through September on compliance inspections. Audits will cover Article 11 technical files for systems deployed after August 2. Providers must present an up-to-date Technical

Documentation dossier before placing a system on the market or putting it into service.

French regulator CNIL, German BfDI and Spanish AESIA will focus initial requests on three sectors. The sectors are automated resume screening tools in human resources, algorithmic credit assessment systems in retail banking and AI triaging tools in private healthcare clinics. Providers of GPAI foundation models that exceed the 10²⁵ FLOPs training threshold will submit systemic risk evaluations by September 15. The European AI Office will review redteaming methodologies, energy consumption disclosures and compliance with the standardized copyright training summary template published in July.

California Senate Bill 1047 cleared both chambers of the legislature in August. Governor Gavin Newsom must sign or veto the legislation by midnight on September 30. If enacted, the bill would impose obligations on developers training AI models that require more than 10²⁶ integer or floating-point operations and cost over \$100 million to train. Colorado will hold rulemaking hearings to finalize reporting forms for Senate Bill 24-205. That law mandates risk management programs to prevent algorithmic discrimination in employment, education, housing and financial services.

Illinois updates to the Illinois Human Rights Act take direct effect.

MEDIA LAB

AI Improves Fake News Detection Then Weakens It

MIT Media Lab researchers tracked 67 participants over four weeks.

Researchers at the MIT Media Lab tracked 67 participants over four weeks as they evaluated news headlines and images and said whether they believed each item was real or fake, sometimes with the assistance of an AI chatbot. Participants were 21 percent more accurate in finding fake news with AI help. Independent discernment did not last after the help was removed. Reliance improved immediate answers and weakened solo judgment.

Some evaluations were completed alone. Some used chatbot support. Assisted sessions produced higher accuracy on the items tested. Unassisted sessions after four weeks showed reduced ability to identify misinformation without the system. The paper describes the outcome as reduced beliefs in

INDUSTRY

Planted Solar, LawZero, Anthropic Secure Compute

\$31.8M solar round, up to \$150M grants each from Canada and Germany, and a Queensland lease tie cash to buildout.

Planted Solar has raised \$31.8M with Google to deploy autonomous solar farms for AI data centers. Canada and Germany have each committed up to \$150M in grants to LawZero. Anthropic has signed its first Australian data-center agreement in Queensland. The three moves tie funding to solar, hiring and capacity. Each attaches money to strict physical requirements for compute. The pattern is concrete. Panels, payrolls and leases come before models. Cash meets ground.

LawZero is the Montreal-based non-profit founded by Yoshua Bengio in 2023 with about \$30M in philanthropic support. The new grants were unveiled at Montreal's All In AI conference. The funding will underwrite hiring and compute for Scientist AI. Planted Solar pairs Google capital with autonomous solar farms built to serve data centers. The structure links generation directly to load. The work is practical. Staff, machines and power draw the schedule. Output is measured in watts and hires. Steel arrives first. Then code.

AUDITS

The EU AI Office begins high-risk audits in September 2026.

BRAZIL

Brazil votes on Bill 2338 in September 2026.

FRONTIER MODELS

GPT-6 Astra triggers critical threshold amid launch rush

September brings five frontier releases with gated cyber tiers and sharp price cuts.



FILE PHOTO Rows of server racks with blinking lights fill a large data center. Capacity on this scale underpins the rapid September model releases and expanding cyber evaluations.
CARL LENDER FROM SUNRISE, USA · CC BY 2.0

Five frontier launches landed in ten days to September 10, with OpenAI GPT-6 Astra triggering a critical cybersecurity threshold while vendors paired releases with steep price cuts and gated cyber variants. Anthropic opened September 1 with Claude Fable 5.1 and Mythos 5.1. Google DeepMind and Meta followed September 2 with Gemini 3.8 Flash and Muse Spark 1.3. OpenAI released Astra September 3 after a September 1 announcement. DeepSeek closed the final run September 10 with V4.1-Flash.

Astra marks the start of the GPT-6 generation after the GPT-5 line. OpenAI placed the production version in limited preview September 3 with rollout to paid and Pro and Enterprise users September 4. The company describes it as its most capable model for work, coding, research and computer use. It retains restricted cyber capabilities. Disclosures position it as the first model to meet the critical cybersecurity threshold, a milestone that frames access controls across the month.

Four of the five launches pair a general model with a

gated security tier. Mythos 5.1 remains behind verification programs while Fable 5.1 is generally available with cache-read pricing cut 75 percent. Gemini 3.8 Flash lists intro pricing through December 31 before doubling January 1. Its Cyber variant is gated behind the Fairwind Program. Muse Spark 1.3 holds a blended price near 10 cents per million tokens, the cheapest model in the ranked leaderboard top five.

“It’s only the first day of September 2026, but the month and fall season are already off to the races in AI land,”

VENTUREBEAT

Capability gains center on efficiency rather than new base architectures. Gemini Terminal-Bench 2.1 scores rose from 81.6 to 90.8, the fourth Flash release since May 19. Muse Spark 1.3 uses about 20 percent fewer tool calls and 25 percent fewer tokens than version 1.2.

MEMORY

Chatbot Study Aid Speeds Work But Weakens Recall

Traditional learners answered 68.5 percent correctly compared with 57.5 percent for chatbot users on a test given 45 days later.

Undergraduate business students who studied with ChatGPT finished faster and remembered less 45 days later. The result comes from a Barcaui experiment with 120 students split into equal groups and tested with a 20 question exam. Traditional study produced stronger retention than chatbot assisted preparation on basic artificial intelligence concepts including ethics, societal impacts and technical foundations. The pattern points to shallower encoding despite faster completion for material requiring conceptual understanding.

Sixty students used the chatbot to explain concepts, generate examples and structure a ten minute presentation. Sixty others used course notes, academic databases and standard internet search engines, with chatbots barred. On the surprise retention test, the traditional group averaged 68.5 percent correct against 57.5 percent for the chatbot group. Eighty five students returned for the twenty question multiple choice exam focused on conceptual understanding rather than memorized definitions alone.

The chatbot group averaged 3.2 hours of preparation against 5.8 hours for traditional learners. The gap in scores remained after adjustment for study time, indicating a difference in quality of study rather than time alone. Traditional work involved re reading materials, self testing and manually connecting ideas, practices that exercise retrieval networks. Chatbot work centered on writing prompts and reading generated responses, activity that feels productive while providing less friction.

“I feel there has been a growing loss of cognition and interest among students over the years.”

ANDRÉ BARCAUI, FUNDAÇÃO GETULIO VARGAS

Loss was largest on highly technical material covering technical foundations. Ethics and societal impacts showed a smaller gap between groups.

CACHE

Anthropic cut cache-read pricing by 75% on September 1.

RED LINES

Ressa Launches Global AI Red Lines Call

The September 2025 declaration signed by 200 politicians and scientists including 10 Nobel laureates urges bans on bioweapon design, mass surveillance and impersonation.

The global call for AI red lines was published online as an open letter on 22 September 2025 and announced by Nobel Peace Prize laureate Maria Ressa at the 80th United Nations General Assembly high-level week in New York. The declaration was initially signed by 200 prominent politicians and scientists, including 10 Nobel Prize winners. The call does not specify which red lines to set, but suggests bans on bioweapon design, mass surveillance and AI impersonation.

The initiative was organized by nonprofit organizations, the French Center for AI Safety, The Future

Society and the Center for Human-Compatible Artificial Intelligence. Ressa urged governments to “define what AI should never be allowed to do” and “establish clear international boundaries to prevent universally unacceptable risks for A.I.” The letter cites risks such as engineered pandemics, widespread disinformation, large-scale manipulation, unemployment and loss of control. It states “an international agreement on clear and verifiable red lines is necessary” and urges governments to reach agreement by the end of 2026 with robust enforcement mechanisms.

HUMANNESS

Message Humanness Predicts Perceived Human Identity

Secondary analysis of 92 women in HeartBot trial links humanness score to misidentification after adjustment.

Secondary analysis of the HeartBot trial found message humanness predicted whether women perceived an AI chatbot as human after adjustment for demographics and prior chatbot use. Among 92 women aged 25 years or older, 61 correctly identified HeartBot as an artificial agent and 31 misidentified it as human. Higher humanness scores were associated with higher odds of human perception, with an adjusted odds ratio of 2.37.

Multivariable logistic regression adjusted for age, race or ethnicity, education, previous AI chatbot use, message humanness, message effectiveness, and attitude toward AI. Only message humanness reached

significance, with a 95 percent confidence interval of 1.26 to 4.48 and P equal to .007. Participants were recruited through social media from October 2023 to January 2024 and interacted with HeartBot once through SMS text messaging.

At the start of the conversation HeartBot introduced itself as HeartBot and did not state that participants would interact with an AI system. Perception of chatbot identity, conversation length, message humanness, message effectiveness, and attitude toward AI were measured in a postchatbot survey. Mean age was 45.9 years with a standard deviation of 11.9 and a range of 26 to 70.

COGNITION

Stimulus Computable Models Test Human Cognition

Frank and Goodman argue human-scale training and experimental probes link AI behavior to theories of human intelligence.

Michael C. Frank and Noah D. Goodman argue that artificial intelligence models can function as cognitive models. The claim appears in their review in Annual Review of Psychology. The central condition is stimulus computability. The models operate over stimuli similar to those experienced by people. Training uses human scale input data. Evaluation uses careful experimental probes. Such models can explain and predict behavior. The authors describe this use as the most exciting opportunity for understanding the human mind.

AI models have a long history as theoretical artifacts in cognitive science. The

current generation differs in one key respect. It accepts inputs that resemble human experience. This property allows direct comparison between model responses and human responses. Researchers can present the same stimuli to systems and to participants. Differences and matches become measurable. The review presents this development as an advance that creates important opportunities for study of the mind. It moves discussion from analogy to test.

The proposal sets two requirements for cognitive modeling. First, training input must remain at human scale.

SEPTEMBER MODEL LEDGER		
VENDOR ANNOUNCEMENT DATES WITH MAKERS AS LISTED IN THE SEPTEMBER TRACKER.		
MODEL	DATE	MAKER
Claude Fable 5.1	September 1	Anthropic
Mythos 5.1	September 1	Anthropic
Gemini 3.8 Flash	September 2	Google
Muse Spark 1.3	September 2	Meta
GPT-6 Astra	September 3	OpenAI

DATED RELEASES	
SEPTEMBER 1	Claude Fable 5.1 ships Anthropic keeps list price unchanged
SEPTEMBER 1	Astra announced OpenAI precedes September 3 release
SEPTEMBER 2	Gemini 3.8 Flash ships Google adds Fairwind-gated Cyber variant
SEPTEMBER 2	Muse Spark 1.3 ships Meta lists contributor tier same evening
SEPTEMBER 3	GPT-6 Astra releases OpenAI starts GPT-6 generation

ASTRA GENERATION	
OpenAI started the GPT-6 generation on September 3. Astra was announced on September 1 and released as GPT-6 Astra on September 3. It was the first model to trigger its critical-cyber safeguard threshold.	