

CYBORG NEWS

CYBORGNEWS.MIT.EDU

TUESDAY, SEPTEMBER 22, 2026 · VOL. II

FREE TO HUMANS

PHYSAI-BENCH - 10,178 instances	UAV decisions	GRICEA REPLAY ▲ 93% replicated	CUI 2026 papers	GRICEA GAPS ▼ 96% missing info	CUI 2026 papers	CODING STUDY - 59 students ChatGPT-4.5 test	RETENTION - 55 retained	recall at 48 hours	MUSE VS CHATGPT ▲ 12 days US Canada launch	SOCIOVERSE 2 ▲ extends 1.0 social simulation
------------------------------------	---------------	-----------------------------------	-----------------	-----------------------------------	-----------------	--	----------------------------	--------------------	---	---

What's News

In brief

- Peer-support tuning uneven SFT and DPO on COPES lifted strategy alignment 50% but gains varied sharply by subreddit and coping strategy.
- MIT Reads turns fictional The 10-year MIT Libraries program pivots to fiction and memoir to rebuild shared humanity in the AI age.
- California AI order weighed EFF backed Newsom's SB-53 loss-of-control reporting push but urged action on bias, Flock surveillance and pricing.
- Health myths countered Culturally adaptive AI videos durably reduced belief in women's health misinformation.

Also noted

- Conversational studies replicated Gricea reproduced 93% of eligible CUI 2026 configurations while finding 96% omitted details needed for replication.
- Society simulated longitudinally SocioVerse2 adds evolving environments, counterfactual forks and editable study states to human-AI simulation.
- Robot identity persists A framework ties continued agent identity to embodied, behavioral and shared-history cues across updates and migrations.
- Digital license assessed IDDL is positioned as demonstrated capability evidence beyond what DigComp and UNESCO define as competence.
- Data quality hospitalized The Data Hospital proposes a ten-stage admit-to-replay workflow separating raw observation from approved interventions.
- Agents drift to collusion Long-horizon LLM agent interactions showed emergent collusive strategies over extended episodes.

Unsealed filings quote Microsoft scientist warning news theft

New York Times-led plaintiffs cite internal ChatGPT and Copilot documents in summary judgment bid.



UNDER REVIEW Stone steps and columns frame the entrance to a federal courthouse. The unsealed filings ask the court to rule on verbatim copying claims involving ChatGPT and Copilot.

W_LEMAY · CC BY-SA 2.0

Unsealed summary judgment filings from news organizations led by The New York Times quote Microsoft Director of Applied Science Brent Hecht calling scraping of news for AI training “an astonishing theft of unprecedented proportions” and perhaps the “largest theft of labor in human history” as Microsoft and OpenAI defend ChatGPT and Copilot against copyright claims.

ChatGPT head Nick Turley warned internally that publishers faced an

“existential threat” from products trained on news that substitute for providers, while a Microsoft document described a “doom loop” that will hurt model performance and the web together. Microsoft data recorded 83 to 93 percent drops in click-through rates for some plaintiffs and 51 to 94 percent drops for others, with low click-through also reported from ChatGPT search results.

Under oath Microsoft CEO Satya Nadella said companies should not

violate terms of use by dodging paywalls, yet OpenAI President Greg Brockman replied “Ah, nice.” after staffer Nick Ryder reported finding “a hack” for crawlers “to get around” the NYT paywall. Nadella also described chatbots stealing clicks by giving information on the AI platform rather than sending users to sources, and Turley agreed there is “no good reason to click” when answers appear directly.

“largest theft of labor in human history”

BRENT HECHT, MICROSOFT

News plaintiffs say they are now ready for trial on “compelling evidence of substitution” and seek a ruling only on articles with “extensive verbatim overlap” from summaries, bullet points and bias ratings.

STARTUPS

a16z launches \$42 million tech academy

The free one-year program places 50 students with partners including Anthropic, Google, Meta and Palantir.

Andreessen Horowitz is launching the Horowitz Andreessen Academy, a free one-year program in San Francisco for 50 recent high school graduates. The project has \$42 million in funding led by a16z. Students will take short classes with founders and operators, including OpenAI CEO Sam Altman, and complete company stints instead of earning a degree.

The academy will launch with 10 partners, including Anduril, Anthropic, Coinbase, Google, Meta, NVIDIA, OpenAI, Palantir, Replit and Stripe. Students will move to San Francisco and arrange housing. Andreessen Horowitz will provide more than \$50,000 in compute credits plus a \$5,000 travel and research budget. There are no traditional grades, tests or homework. Time is split between instruction, work and exploration.

“political lobbying operations fighting American innovation”

MARC ANDREESSEN, ANDREESSEN HOROWITZ

LEARNING

ChatGPT Raises Coding Scores But Lowers Student Recall

A 55-student trial in C programming finds higher task scores paired with lower recall and ownership than web search.

Students using ChatGPT-4.5 scored 89 percent on three introductory C programming tasks, compared with 69 percent for students using conventional web search without generative AI. The same ChatGPT group recalled less of the material immediately and after 48 hours, and attributed less of the submitted code to themselves. The controlled between-subjects experiment retained 55 undergraduate computer science students for analysis.

Recall was measured by cued recall tests given immediately and 48 hours later. ChatGPT-assisted students scored 41 percent versus 53 percent on the immediate test. After 48 hours, the scores were 39 percent versus 52 percent.

The study found no significant difference in the loss of recall information over 48 hours between the groups.

Ownership showed a sharp gap. Students with ChatGPT attributed 45 percent of the submitted code to themselves, compared with 81 percent in the web search group. Self-reported mental effort increased less across tasks in the ChatGPT condition, with a Holm-adjusted p value of .047. The study also measured self-reported difficulty.

The study describes a controlled between-subjects experiment with 59 undergraduate computer science students, with 55 retained for analysis. Participants completed three introductory C programming tasks with

access to ChatGPT-4.5 or conventional web search without generative AI. Measures included task performance, self-reported mental effort and difficulty, pupillary responses, heart rate variability, and ownership.

Physiological tracking did not change the picture. Confirmatory tests of pupillary responses and heart rate variability did not detect significant differences in trajectories between conditions. The study notes substantial data loss, which limits interpretation of those signals. The core pattern remained in the behavioral measures. Assisted performance rose while recall and ownership fell. The gap appeared immediately and persisted at 48 hours.

APPS

Muse Outpaces ChatGPT in Early iOS Downloads

Apptopia estimates 1.8 million U.S. and Canada downloads in the first 12 days.

Market intelligence provider Apptopia estimates Meta Muse passed ChatGPT in early iOS downloads in the U.S. and Canada. Muse recorded 1.8 million downloads in its first 12 days. ChatGPT recorded 1.3 million in its first 12 days. The comparison covers iOS in those two markets only. Muse also leads in daily active users.

Apptopia aligned the figures to account for different launch strategies. ChatGPT launched on iOS globally. Muse launched on the App Store and Google Play in the U.S. and Canada only. The firm compared iOS activity in the U.S. and Canada. Total Muse installs reached 2.8 million globally in the first 12 days.

Muse holds 642,000 daily active users in the U.S. on mobile. ChatGPT held

231,000 at the same point after launch. On iOS alone Muse holds 359,000 daily active users. The total remains above the ChatGPT figure from that period. Muse moved from No. 2 overall on the U.S. App Store after launch to No. 1. Business Insider reported the rise on Friday. Appfigures said Muse had crossed 1 million downloads.

Apptopia provides third party estimates only. The firm lacks direct access to Meta internal figures. Meta has not shared public figures related to Muse early adoption. Meta used cross promotion for Threads on Instagram and Facebook. Threads now counts more than 500 million users. Muse connects with both platforms and works inside WhatsApp.

PRIVACY

Default Chatbot Settings Retain Private Conversations For Training

Enterprise deletion contracts remain out of reach for most people, leaving policy promises and cryptographic tools as the available options.

ChatGPT, Claude and Gemini now receive the material once reserved for therapists and confession booths. The three services collect and store those exchanges by default. Stored text can be used for further training, shared with partners and contractors, and produced for plaintiffs and law enforcement with a legal order. Practical protection requires changed settings, different products, or enterprise contracts.

The baseline for a standard account is near zero privacy against a determined holder of legal process. The service owner can read the logs. Advertisers and partner companies can receive them. Contractors who fine tune systems can review them. Police and civil litigants can demand them. Enterprise buyers can purchase stronger contractual terms.

That stronger term is zero data retention, known as ZDR. OpenAI, Anthropic and Google offer it for paid enterprise and developer accounts. Records are

deleted immediately after processing when it is enabled. It is not offered to average users. It contains exceptions for abuse monitoring. Those exceptions still permit analysis and logging before deletion.

“People are integrating AI into their personal lives. They talk with it about their deepest insecurities, their finances, their health, their relationships.”

MOXIE MARLINSPIKE, CONFER

Anthropic does not offer ZDR for its most sophisticated Mythos class models like Fable 5.1. The stated reason is potential misuse including scamming, hacking and autonomous misbehavior. OpenAI analyzes user activity on customer systems before deletion under Private Safety Processing. It flags abuse for the customer organization

and can alert OpenAI staff without revealing content. Google logs some Gemini prompts for abuse monitoring without Google ID or IP address. Sanitized prompts can still identify a person when the text itself contains identifying detail.

Most individuals must rely on weaker policy promises or separate architectures. Proton offers Lumo as AI where every conversation is private. The protection is a company promise not to log, not end to end encryption. Venice.ai and Duck.ai make similar no log vows while relaying requests to models including Claude and ChatGPT. Duck.ai lets the user select the underlying model. Venice.ai routes requests automatically without always disclosing the destination. Confer uses cryptography to prevent its own server from surveilling or logging conversations.

Cryptographer Moxie Marlinspike created Signal in 2014 after concern about unprotected texting. A decade ago text messages held the most sensitive personal data from devices.

YOUTH SAFETY

EU Proposes Phased Social Media Access For Minors

Draft law ties access to age assurance and parental controls, drawing privacy warnings from the Electronic Frontier Foundation.

The EU Commission has proposed age phased access to social media and video sharing platforms backed by broad age assurance checks. The draft law reserves full access for adults and places younger users in tiered accounts with parental controls. The Electronic Frontier Foundation warns the plan puts services behind age gates and expands intrusive verification for all users.

Under the proposal platforms deemed risky face a phased access schedule. Children under 13 would have no service accounts. Users from 13 to 15 would hold restricted accounts under tight parental supervision. Users from 15 to 18 would hold autonomous accounts in a safe by design environment. The threshold for risky covers mainstream services that rely on personalized recommender systems or offer uninterrupted content consumption.

The delay is tied to system wide age verification based on the proposed EU age verification scheme.

Parents would create accounts for teenagers and prove parental status in a separate check. Existing accounts could avoid a new check only where a provider can show with a high degree of confidence that the user is above the relevant threshold. The Foundation describes that standard as vague and the process as a compliance burden with privacy costs.

“However they’re designed, age gates undermine civil liberties, reduce safety, and create barriers to internet entry, often at the expense of marginalized groups.”

CHRISTOPH SCHMON,
ELECTRONIC FRONTIER
FOUNDATION

Safety by design rules would reach beyond social media to video sharing services online games AI companions chatbots and app stores.

SCALE AGAINST RESULT	
WHAT EACH EFFORT MEASURES WHAT IT FOUND.	
SYSTEM	SCALE
PhysAI-Bench	10,178 instances
Gricea rebuild	93% of eligible papers
Gricea audit	96% of papers
Coding study	59 students enrolled
Coding study	55 retained
ChatGPT-4.5 arm	higher coding score
COPES	Reddit peer queries
COPES eval	three-axis framework
SocioVerse2	extends SocioVerse 1.0
SocioVerse2	co-evolutionary paradigm
Muse mobile	first 12 days
Muse mobile	US and Canada only
MIT Reads	10-year program
Data Hospital	approved interventions only
EU Kids Act	age gates
Embodied AI	updates and migration

METHODS

MIT Methodologist Tests How Far Social Science Findings Travel

Naoki Egami studies external validity, AI measurement and campaign field experiments.

Naoki Egami is an MIT political scientist who studies whether research results apply beyond the settings that produced them. He specializes in methodology and in external validity. His portfolio spans campaign field experiments and studies of artificial intelligence tools in research. He joined the Department of Political Science as an associate professor with tenure in 2025. He is also a faculty affiliate of the Statistics and Data Science Center.

“You need both perspectives,” Egami says. He describes political methodology as asking questions as a political scientist and solving them as an applied statistician. He says researchers need technical theory and attention to empirical problems. Without theory, solutions stay ad hoc. Without problems, theory misses substance. His approach seeks optimal solutions to mathematical problems behind empirical challenges. “You find out the underlying mathematical problems

behind the empirical challenges people face, and solve them optimally.”

“You find out the underlying mathematical problems behind the empirical challenges people face, and solve them optimally.”

NAOKI EGAMI, MIT

In statistics, similarity often means population data. In political science, context differs. Egami examines campaign research to show the gap. Researchers sometimes receive permission to run field experiments in active campaigns. Politicians grant access when they expect triumph. Voter behavior in those races can differ from behavior in competitive districts. Findings from those races may not transfer to battlefield races.

Before the current artificial intelligence interest, Egami began studying AI tools in studies.

DRONES

Benchmark tests 29 models on drone decisions

Top accuracy reaches 52.00 percent on 500 network-constrained decisions.



FIELD TEST A small quadcopter holds position above an open grass field. PhysAI-Bench measures whether models can make such flight decisions under latency, packet loss and edge load.

U.S. AIR FORCE PHOTO BY SENIOR AIRMAN AMANDA JETT · PUBLIC DOMAIN

PhysAI-Bench evaluates 29 foundation models on 10,178 standardized decisions from autonomous UAV missions. Each decision includes mission context, sensor observations, physical constraints and AI-native 6G network conditions. Conditions include latency, packet loss, throughput, edge load and network slicing. GPT-5.3 records the highest accuracy at 52.00 percent. The result leaves reliable agentic decision-making unresolved.

The benchmark exposes only information preceding each decision. The procedure prevents future-event leakage and approximates online decision-making. Evaluation follows a two-

stage protocol. First, model-specific configurations are selected from 12 combinations of zero, three and five-shot prompting and four temperatures. Selection uses three runs on a 35-instance human-verified development set. Then each selected configuration is frozen and tested in three runs on a fixed episode-disjoint set of 500 instances.

GPT-5.3 achieves 52.00 percent accuracy. GPT-5.2 follows at 49.40 percent. Grok 4.5 follows at 49.07 percent. Few-shot prompting generally improves performance. Temperature has limited influence overall.

AGENTS, SUPPORT AND REPLICATION

Social simulation now tests intervention. SocioVerse2 extends SocioVerse 1.0 into a human-AI co-evolutionary paradigm built to give researchers control over content and over process. Earlier platforms verified collective behavior and aligned simulated populations with real societies in cross-sections and used autonomous agents for research. Two requirements remained without systematic support: intervention in the content of a simulation and researcher control over the process that produces it. The new framework addresses both.

Peer support on Reddit often goes unanswered, and clinical benchmarks do not measure community alignment. The COMMunity-centered Peer Engaged Support dataset addresses lived-experience informed support for mental health queries. A three-axis evaluation framework assesses alignment with community perspectives. Large language models have shown strong performance on clinical benchmarks, but generation of lived-experience informed and community-aligned peer support remains underexplored. The work tests whether fine tuning brings heterogeneous benefits for support generation.

Conversational AI research remains hard to replicate because systems and study configurations are reported in fragments. Gricea represents studies as configurable and deployable research artifacts that researchers can run and inspect and share and reuse. In a replication study the platform rebuilt configurations for 93 percent of eligible CUI 2026 papers while flagging missing information in 96 percent. Physical AI faces a parallel evaluation gap for agentic decision making in dynamic environments. PhysAI-Bench contains 10,178 standardized decision instances extracted from conversational traces of autonomous UAV missions — preserving mission context and temporal dependencies and physical constraints. In programming education, 59 undergraduate computer science students completed three introductory C tasks with ChatGPT-4.5 or conventional web search, with 55 retained for analysis. ChatGPT-assisted students achieved higher coding scores, with cued recall measured immediately and 48 hours later alongside mental effort and difficulty and pupillary responses and heart rate variability and ownership.