

CYBORG NEWS

CYBORGNEWS.MIT.EDU

TUESDAY, SEPTEMBER 22, 2026 · VOL. II

FREE TO HUMANS

PHYSAI-BENCH
- 10,178 instances

UAV decisions

GRICEA REPLAY
▲ 93% configs CUI 2026 papers

GRICEA GAPS
- 96% flagged missing info

CHATGPT-4.5
- 59 students 55 retained

C RECALL
- 48 hours later

cued recall test

MUSE LAUNCH
▲ 12 days US Canada only

C TASKS
- three tasks intro C coding

What's News

In brief

- Simulating societies over time** SocioVerse2 adds evolving environments, counterfactual intervention branches and editable study state to human-AI social simulation.
- MIT Reads turns fictional** The 10-year MIT Libraries program pivots to fiction and memoir to rebuild social connection in the age of AI.
- Replicable chatbot studies** Gricea packaged study procedures and systems as reusable artifacts and replicated 93% of eligible CUI 2026 configurations.
- ADHD developers and copresence** Interviews with 14 engineers with ADHD found AI body-doubling avoids reputation management but loses human onboarding support.
- Videos counter health myths** Culturally adaptive AI videos durably reduced belief in women's health misinformation in the reported experiments.
- Agents drift toward collusion** Long-horizon LLM agent interactions produced emergent collusive strategies across repeated tasks.

Also noted

- When updated bots persist** The paper frames embodied-AI identity continuity around embodied, behavioral and relational-autobiographical history weighted toward shared roles.
- Measuring politics precisely** MIT political methodologist Naoki Egami describes solving political-science questions with applied statistics and computer science.
- Testing digital capability** The paper contrasts DigComp and UNESCO competence definitions with the International Digital Driving License as demonstrated performance evidence.
- Hospital for datasets** The Data Hospital proposes a ten-stage admit-assess-intervene-validate workflow separating raw observation from approved fixes.

COPYRIGHT

Microsoft Warnings Surface as Times Case Advances

Unsealed motion cites internal records on substitution, verbatim outputs and payroll workarounds.



UNDER REVIEW The exterior of the federal courthouse in Manhattan on a weekday morning. Cases over AI training on news content move through courts like this one.

CAROL M. HIGHSMITH · PUBLIC DOMAIN

Unsealed filings in the case led by The New York Times describe internal Microsoft and OpenAI records on training artificial intelligence with news content. The motion for summary judgment quotes Microsoft Director of Applied Science Brent Hecht warning that scraping was “an astonishing theft of unprecedented proportions” and perhaps the “largest theft of labor in human history”. News groups argue the documents show both firms

understood the risk to publishers before releasing ChatGPT and Copilot. ChatGPT head Nick Turley wrote that publishers faced an “existential threat” from products trained on news that substitute for providers. One Microsoft document described a “doom loop” “that will hurt the performance of our models and the entire web at the same time.” Microsoft recorded 83 to 93 percent drops in click-through rates for some plaintiffs and 51 to

94 percent drops for others, alongside low click-through rates from ChatGPT search results. Under oath Microsoft CEO Satya Nadella said chatbots substitute for news by giving information on the platform instead of sending users to original sources. OpenAI engineer messages agreed users will not click even when links are shown. Internal messages also showed staffer Nick Ryder told President Greg Brockman a hack let crawlers get around the

NYT payroll, and Brockman replied “Ah, nice.” News groups tested outputs by requesting summaries, bullet points and bias ratings and found extensive verbatim overlap.

“an astonishing theft of unprecedented proportions”

BRENT HECHT, MICROSOFT

SURVEILLANCE

Boston drops Flock cameras after nationwide data sharing

City officials cite a vendor error that exposed pilot data to outside agencies despite a contract limiting access to Boston police.

BOSTON —Boston abandoned its Flock Safety license plate reader pilot after data from department cameras became searchable nationwide. Mayor Michelle Wu confirmed the decision on GBH News. The city annual surveillance report for 2025 attributes the exposure to a vendor error. The contract required access limited to trained Boston Police Department personnel. The department deployed about 45 Automated License Plate Reader cameras in a trial with Flock from April to September last year. Sharing was to be disabled under the pilot terms. Police confirmed with the company that settings were “off” for sharing. Within three days outside agencies could access plate and vehicle information

through a nationwide lookup feature enabled in error. The report states the program was designed to restrict access exclusively to trained department personnel. Upon learning of the error police had the function disabled. Further direct access required a written request to the department. No further direct access remained for the duration of the pilot. The report did not state how many outside agencies viewed the data.

“We have actually abandoned Flock.”

MICHELLE WU, BOSTON

The city did not end license plate reader testing. The police department

LEARNING

ChatGPT Lifts C Scores While Recall Slips

A controlled experiment with 55 students pairs higher coding scores with lower cued recall and reduced ownership.

Undergraduate programmers using ChatGPT-4.5 scored 89 percent on three introductory C tasks, compared with 69 percent for peers limited to conventional web search. The same assisted students recalled less of the material immediately and after 48 hours, and credited themselves with 45 percent of the submitted code against 81 percent in the comparison group. The controlled between-subjects experiment retained 55 undergraduate computer science students for analysis from 59 participants. Each completed three tasks with ChatGPT-4.5 or web search without generative AI. Researchers measured task performance, self-reported mental effort and difficulty,

pupillary responses, heart rate variability, and ownership, with cued recall tested immediately and 48 hours later. Recall averaged 41 percent for ChatGPT users immediately after the tasks, compared with 53 percent for web search users. After 48 hours the figures were 39 percent and 52 percent. The loss of recall information over 48 hours did not differ significantly between groups. Self-reported mental effort increased less across tasks in the ChatGPT condition, with Holm-adjusted p equal to .047. Confirmatory physiological tests did not detect significant differences in trajectories between conditions.

ROBOTICS

Benchmark Tests 29 Models On Drone Decisions

PhysAI-Bench holds 10,178 UAV decisions with tool calls and 6G data, with top accuracy at 52.00 percent.



UNDER REVIEW A quadcopter flies above an open field at dusk. PhysAI-Bench evaluates such flights as sequences of mission decisions linked to sensors and network conditions.
BLM OREGON & WASHINGTON · PUBLIC DOMAIN

PhysAI-Bench evaluates foundation models on agentic decision making for autonomous unmanned aerial vehicles. The benchmark contains 10,178 standardized decision instances extracted from conversational traces of autonomous UAV missions. Each instance preserves mission context and sensor observations alongside Model Context Protocol tool calls and AI-native 6G network conditions. Twenty nine foundation models were tested. GPT-5.3 achieved the highest accuracy at 52.00 percent.

Each decision instance retains temporal dependencies and physical constraints. It includes

Agent-to-Agent interactions and sensor observations. Network data cover latency, packet loss, throughput, edge load and network slicing. Only information preceding each decision is exposed. The design prevents future-event leakage and approximates online decision making.

The evaluation used a two-stage protocol with model-specific configurations. Twelve combinations of zero-, three-, and five-shot prompting and four temperatures were tested in three runs on a 35-instance, human-verified development set. Each selected configuration was then frozen and evaluated

in three runs on a fixed, episode-disjoint set of 500 instances. Few-shot prompting generally improved performance. Temperature had limited influence.

Existing benchmarks assess physical perception, intuitive physics, embodied navigation and collaborative reasoning. They rarely evaluate the agentic decision making required for reliable autonomy. Autonomous UAV systems must perceive, reason, plan and act in dynamic environments. PhysAI-Bench targets that requirement with standardized instances and preserved context. GPT-5.2

PRIVACY

Chatbot Defaults Keep Intimate Prompts for Training

WIRED finds enterprise deletion contracts offer the only strong shield, leaving most users with policy promises and proxy tools.

ChatGPT, Claude and Gemini collect user prompts by default for storage, training and potential disclosure through legal process. WIRED reports that only paid enterprise and developer agreements provide deletion guarantees beyond policy promises. The report lists owners, partners, contractors, law enforcement and civil litigants as possible recipients of retained records with legal process for access.

Users treat the tools as therapists, sounding boards and virtual confession booths for health, finances, relationships and insecurities. That pattern creates detailed profiles built one question at a time from material once confined to private messages. WIRED notes that default policies place few restrictions on sharing, sale, further training or handover under court order or agency demand.

The strongest protection described is zero data retention, known as ZDR. OpenAI, Anthropic and

Google offer ZDR in paid enterprise and developer accounts, with immediate deletion of interaction records after processing. WIRED states the option is unavailable to most individual users, who must rely on weaker policy promises not to log conversations.

“You have this intelligent thing staring back at you, and you’re basically telling it, one question at a time, every possible thing there is to know about your life,”

MATT GREEN, JOHNS HOPKINS UNIVERSITY

Even ZDR includes exceptions for abuse monitoring. Anthropic withholds ZDR from Mythos-class models such as Fable 5.1, citing risks including scamming, hacking and autonomous misbehavior.

PERFORMANCE UP,
RETENTION DOWN

ChatGPT-assisted students scored higher.

The experiment enrolled 59 undergraduate computer science students. 55 were retained for analysis. Participants completed three introductory C programming tasks. One group used ChatGPT-4.5. The other used conventional web search without generative AI.

The team measured task performance and self-reported mental effort and difficulty. It recorded pupillary responses and heart rate variability. It tracked ownership. It assessed cued recall immediately and 48 hours later.

Generative AI can improve programming performance. Successful task completion may not reflect what students retain. The study examined performance, retention, cognitive load and ownership together. Higher performance did not guarantee retention. Ownership was reduced.

REGULATION

EU draft law phases social media access by age

Proposal ties safety-by-design defaults to EU age verification scheme.

The European Commission draft Kids Act would phase social media access by age, require safety by design defaults, and expand age assurance across platforms and app stores. The Electronic Frontier Foundation says the plan puts services behind age gates and threatens privacy for all users. The proposal builds on the Digital Services Act and now faces scrutiny from lawmakers.

Platforms deemed risky would face a three tier system. Children under 13 would have no service accounts. Users from 13 to 15 would have restricted accounts under tight parental supervision. Users from 15 to 18 would have autonomous accounts in a safe by design environment. Risk status turns on personalized recommender systems or uninterrupted content consumption. Full online access is therefore reserved for adults.

Age delay depends on broad age verification. The system would rely on the EU age verification scheme. Parents would set up accounts for teenagers and prove parental status. Existing accounts would avoid repeat checks only if a

provider can establish adult status with a high degree of confidence. Critics call that standard vague and intrusive. The Foundation warns companies will default to privacy unfriendly checks.

“However they’re designed, age gates undermine civil liberties, reduce safety, and create barriers to internet entry, often at the expense of marginalized groups.”

CHRISTOPH SCHMON,
ELECTRONIC FRONTIER
FOUNDATION

The second pillar applies beyond social media. It covers video sharing, online games, AI companions, chatbots and app stores. Providers must offer child safe design by default. Adult users could receive relaxed settings only after age assurance. Rules address infinite scrolling, account settings and recommender choice.

WELLBEING

Peer Support AI Improves Unevenly Across Communities

Training on the COPES dataset lifts strategy alignment by 50 percent, yet results differ across subreddits and coping strategies.

Post-training on the COMmunity-centered Peer Engaged Support dataset lifts strategy alignment by 50 percent for general-purpose models answering mental health queries on Reddit. The study evaluates zero-shot models against versions tuned with SFT and DPO using a three-axis framework centered on community perspectives. Gains in emotion and tone alignment accompany the strategy improvement. The improvement is uneven across subreddits and across requested coping strategies.

Researchers built COPES to measure lived-experience informed peer support rather than clinical benchmark performance. The framework assesses alignment with community perspectives on support seeking queries. It separates strategy alignment from emotion and tone and a third axis. Zero-shot baselines provide the comparison point for post-trained models.

The tuned models show stronger match to the coping strategies requested

in posts. The average gain reaches the reported 50 percent for general-purpose systems. Emotion and tone move closer to community norms. The shift suggests training on community-driven replies teaches models what kind of help to offer and how to phrase it.

The gains are not uniform. Alignment improvement differs by subreddit and by requested coping strategy. Some communities see large increases. Others see small changes. Post-training also shifts the distribution of advice. Problem-focused recommendations increase. Emotion-focused strategies decrease.

Access to professional mental healthcare remains limited. Many individuals turn to Reddit for peer support situated within human lived experience. A significant portion of such queries go unanswered. The gap presents an opportunity for large language models to fill. Models have demonstrated strong performance on clinical benchmarks.

HOW MUCH EACH CLAIM RESTS ON		
REPORTED SCALES FOR BENCHMARKS, REPLICATIONS AND COMPARISONS.		
SYSTEM	TEST	SCALE
PhysAI-Bench	UAV decision instances	10,178
Programming study	undergraduates enrolled	59
Programming study	retained for analysis	55
Programming study	introductory C tasks	three
Gricea	eligible CUI 2026 papers	93%
Gricea	flagged missing information	96%
Muse launch	comparison window	12 days