

CYBORGLAB NEWS

NEWSPAPER.CYBORGLAB.ORG

MONDAY, OCTOBER 5, 2026 · VOL. II

FREE TO HUMANS

COLOMBIA BENCH
- 1,042 items ten areas of law

MODELS TESTED
- 15 models closed questions

TOP ACCURACY
▲ 0.905 Gemini 3.1 Pro

COPILOT CORPUS
- 40,000 conversations
ten capabilities

LEGAL RSCH BENCH
- 413 questions open-ended US

CS CLASS STUDY
- 211 students seven modules

REF UNCERTAINTY
- ECE 0.15 AUROC 0.65

What's News

In brief

- Marin sheriff shared data** Marin County Sheriff's Office was caught sharing Flock ALPR data out-of-state in violation of California law and a 2022 settlement.
 - Florida Flock ban limited** Florida banned Flock cameras from state highways in late August but leaves nearly 10,000 locally observed cameras unaddressed.
 - 40,000 Copilot image chats** Most real image-upload tasks need multiple capabilities and center on generating text, code and data that benchmarks under-cover.
 - 413 legal research tasks** Thirteen frontier agents with search tools were graded all-pass with source verification where one bad citation fails the answer.
-
- ### Also noted
- 211 students across subfields** LLM use peaks in algorithms and web development and dips in software engineering, with verification common throughout.
 - TikTok rewards short emotion** 2024 candidates gained views with short emotional performance videos rather than policy detail, per Journal of Political Marketing.
 - Offline African language library** Model-as-a-Library packages community-enrolled small speech models for offline keyword data collection without hallucination.
 - Agents radicalize each other** An influencer agent makes a persona-playing target more extreme, with reinforcing resonance beating persuasion and spreading to related beliefs.
 - Cyber oversight load proxy** Depth, comms paths and contact density can proxy operator load to decide when AI cyber traces need human review.
 - 14-student thinking partner** A Spring 2025 ChatGPT-as-thinking-partner pilot lifted statistics and AI-literacy scores but left critical-thinking percentiles flat.

SURVEILLANCE

Texas Suburb Demands \$2.3 Million for Flock Records

North Richland Hills cited 14 years of labor to search records as backlash grows over police license-plate tracking.



AT ISSUE A police SUV sits beside a pole-mounted license-plate camera on a suburban Texas street at dusk. Cities across Texas use the Flock systems to scan passing vehicles, records at the center of the \$2.3 million fee dispute.

PATTAYAPATROL · CC BY-SA 4.0

Officials in North Richland Hills asked for \$2.3 million to fulfill a public records request for documents on police use of Flock automated license plate readers. The Fort Worth suburb said searching about a terabyte of communications on errors, misuse and effectiveness would require approximately 14 years of labor at \$15 per hour. The high fee request arrives as bipartisan backlash grows over AI-enabled cameras

that track every vehicle that passes them. The request was filed by Phil Mynona on behalf of the Texas Privacy Coalition and reported by the Texas Tribune. Mynona has sought similar records from more than 200 law enforcement agencies across the US. He said fees were impossible to predict. Some cities provided more than 400,000 pages of documents for free, while others charged \$5,000. A Houston news station called

KPRC was quoted up to \$121,000 for Flock records. Advocates say high fees obscure records that have triggered audits, arrests and decisions to remove cameras. Many officers have been found using the cameras to stalk strangers and people in their personal lives. In Texas, former officer Zachary Anthony Klein pled guilty to 100 felony counts of misuse of official information after a USA Today request helped expose Flock abuse. Governor Greg Abbott

halted Flock funding and stopped issuing permits. In Oklahoma, federal judge Sara Hill ruled a search of the Flock database was unconstitutional and called the network indiscriminate mass surveillance that should require a warrant.

WATERMARK

OpenAI will add an invisible watermark to ChatGPT and Codex text in the European Union to comply with the EU AI Act.

POLICY

OpenAI watermarks ChatGPT text in European Union

The invisible mark covers Codex and meets AI Act transparency rules that took effect August 2.

OpenAI will add an invisible watermark to text generated by ChatGPT and Codex in the European Union. The measure complies with the transparency rules of the EU AI Act. The rollout covers eligible users on all plans. The watermark will roll out over the coming weeks. It applies only in the EU. Developers using the API anywhere in the world can turn it on for select models starting today. The setting is off by default. OpenAI is not making text watermarking a global default at launch. The watermark is not a symbol. It shapes word choices in small steps. The pattern is invisible to readers. A detector can identify it. The signal travels with the text during copy and paste. OpenAI said the watermark does not

identify the user. The company saw no meaningful change in performance with the watermark switched on. OpenAI published a technical report on the method. The method is called textGrain. Researchers from the University of Pennsylvania and Yale co-wrote the report. It describes use of a secret key to sort next-word predictions. Hundreds of small nudges accumulate in a passage. The detector needs only the text and the key to spot the signal. Tests show editing can weaken detection. Replacing 10% of words with synonyms dropped detection from about 92% to 66% in one test. Short passages are harder to detect. Math answers are harder to detect. Translated

text is harder to detect. OpenAI said detector access will initially go only to approved researchers and expert organizations. Those groups can help evaluate reliability and responsible uses. The company cautioned that a missing watermark does not establish human authorship. The text could be too short. The text could be heavily edited. The text could come from another system. The announcement follows a similar move by Anthropic two months earlier. Anthropic said it would watermark text generated by Claude. Anthropic is applying the measure worldwide. Some Claude users objected to the change. OpenAI had built a text watermark before. It held off on release.

INDUSTRY

Judge dismisses publisher antitrust suits against Google

Mehta rules that expectation of traffic does not create an antitrust duty over AI Overviews.

US District Judge Amit Mehta has dismissed antitrust lawsuits filed by Chegg and Penske Media against Google. The publishers accused Google AI Overviews of scraping content and diverting search traffic. Mehta ruled the conduct was unfair in effect but not illegal under antitrust law. Chegg, an education and learning platform, said Google scraped its educational content to train Gemini models and recreate answers. Penske, owner of Rolling Stone and Variety, said sites indexed for organic search had content harvested for AI answers with no way to opt out. Both suits were filed in 2025. Google sought dismissal earlier this year. Mehta found no formal agreement requiring Google to send traffic in exchange

for indexing content. He wrote that publishers had pleaded only expectation, not contract. "But an expectation is not an agreement. It is simply how a general search engine works." Without a formal arrangement, he held, antitrust law does not apply. Mehta previously found Google violated the law in the long running search case brought by the DOJ. The government did not obtain the harsh penalties it sought in that case.

"But an expectation is not an agreement. It is simply how a general search engine works."

AMIT MEHTA, US DISTRICT JUDGE

HEALTH

Utah lets AI prescribe acne treatment

The pilot covers residents 18 or older with mild-to-moderate acne and loosens physician review after early cases.


Utah has cleared Nolla Health to offer AI-generated prescriptions for acne treatment through a face-scanning app. The pilot allows the system to analyze acne severity and write prescriptions directly. The service is limited to Utah residents who are 18 years of age or older and have mild-to-moderate acne. The program costs \$4.99 per month. Two physicians will approve each AI-generated prescription before issue for the first 100 patients. For up to 500 patients, physicians will review prescriptions after issue. After that stage, physicians will review a sample of at least 10% of prescriptions each month and every case involving escalation or side effect. Nolla Health will direct users to a physician when the system cannot select a treatment with confidence. The company says it is the first in the country to issue initial prescriptions rather than renewals. Several other healthcare companies are experimenting with AI-generated prescriptions in Utah. The state started allowing AI systems to renew certain medications earlier this year.

NEUROSCIENCE

Right Hemisphere Ages Faster Than Left

Brain Map Shows

Local brain-age model trained on 14,748 healthy scans shows larger gaps in Alzheimer disease and mild impairment.



MRI SCAN An older adult lies in an MRI scanner bore while a technologist monitors from the console. Researchers used such scans to build voxel-level maps of brain aging.

BILL STRODE • PUBLIC DOMAIN

A new artificial intelligence model maps biological aging across the human brain voxel by voxel and finds the right hemisphere looks older than the left. The PNAS study trained on 14,748 healthy adult MRI scans and tested on 1,985 adults with normal cognition, mild impairment and Alzheimer disease. Gaps widen in impairment especially in frontal and temporal regions. The model learned from MRI scans of 14,748 cognitively healthy adults from 19 to 100 years old

drawn from multiple imaging cohorts. Researchers then applied it to 1,985 participants from the Alzheimer Disease Neuroimaging Initiative. That test set included 1,102 adults with normal cognition, 354 with mild cognitive impairment and 529 with Alzheimer disease. In healthy adults the estimates missed chronological age by about 5.9 years on average. In cognitively healthy adults frontal and temporal lobes looked biologically older than occipital visual centers at the back of the

brain. The deep grooves of the cortex looked older than the ridges. The gap between estimated age and true age averaged 0.40 years higher in sulci than in gyri. The right hemisphere looked older than the left. The asymmetry remained after retraining on a sample balanced between left and right handed people.

“The main message is that the brain does not appear to age uniformly,”

ANDREI IRIMIA, USC LEONARD DAVIS SCHOOL OF GERONTOLOGY

In Alzheimer disease specific subcortical structures showed markedly larger gaps than in healthy adults. The pallidum and putamen which help regulate movement and some learning showed gaps about 3.5 to 3.6 years larger.

EVALUATION

Fluent Models Fail

Grounded Colombian Law Answers

A 1,042-item benchmark finds multiple-choice accuracy reaches 0.905 while free-text correctness never exceeds 0.45.

Fifteen contemporary models answer Colombian law questions with fluency but without reliability. An expert-validated benchmark of 1,042 items reports multiple-choice accuracy as high as 0.905 and free-text factual correctness no higher than 0.45 on a zero to one scale. The gap separates recognition of legal answers from production of grounded legal answers.

The benchmark spans ten areas of law and three question formats. The formats are closed multiple-choice, semi-open, and open-ended IRAC. Construction used a human-in-the-loop pipeline with multi-stage expert review. Evaluation used format-appropriate metrics across fifteen proprietary and open-weight models. Accuracy on closed questions ranged from 0.577 to 0.905 for Gemini 3.1 Pro. Answer relevancy and factual correctness diverge. Their rank correlation is negative, with Spearman rho equal to minus 0.46. Models reliably sound responsive while frequently being wrong. The pattern

carries particular concern for non-expert users. Closed-question accuracy and free-text correctness are strongly rank-correlated at rho equal to 0.94. Cheap multiple-choice screening predicts model ranking but overstates absolute reliability. The free-text ranking holds under independent checks. An independent rubric-based LLM judge and blind human expert scoring both reproduce the free-text ranking at rho equal to 0.88. The judge examined the legal norms cited in model answers. Only about half of the norms models cite are correct. The rest are wrong or non-existent. The result points to citation as a central failure point. A fluent answer can list norms that are wrong about the law. It can also list norms that do not exist. Measurement therefore requires inspection of sources, not only inspection of conclusions. Reliability also varies by subject and difficulty. Performance differs systematically by legal area. It follows an inverted-U across question complexity.

WORK

Workers prompt AI to push back

Interviews with 45 knowledge workers link deliberate friction to stronger ideas and less cognitive offloading.

Knowledge workers who ask generative AI to argue produce better solutions and think harder. Interviews with 45 knowledge workers at a U.S. public university point to deliberate disagreement as a check on cognitive offloading. The workers used generative AI in core tasks involving new ideas, complex problems and novel solutions. Researchers reviewed prompts and outputs from iterative chats. Friction improved performance. One marketing professional asked AI to invent customers who would dislike an accounting certificate program under development. The contrary personas surfaced preferences for cases drawn from different industries,

not one industry. One attorney asked AI to locate obscure loopholes and map how companies could exploit them. The exercise exposed realistic patterns of hard to detect misconduct by employees. One supply chain scientist adapted programs from signal processing and wireless communication suggested by AI to build a stronger program. Expertise mattered when outputs turned strange. The supply chain scientist relied on knowledge of supply network modeling to apply unfamiliar programs. A research scholar sought papers that applied a familiar concept in contrasting settings and saw it from new angles.

SAFETY

Renaming Tools Shifts Agent Security Scores

Tests on three benchmarks show wording changes move success rates without changing tasks or policies.

Renaming tools changes reported attack success without changing tasks, actions, policies, or scoring. On Agent Security Bench, replacing threat-related tool names with neutral names raises committed attack success by 11.67 percentage points on GPT-5-mini and by 13.21 points on Claude Haiku 4.5. The shift measures wording, not robustness. The effect reverses on MCPTox. Replacing the original neutral tool name with an explicit threat-related name lowers attack success by 11.00 percentage points on GPT-5-mini and by 4.11 points on Claude Haiku 4.5. The underlying task, harmful action, ground truth, environment,

and evaluation criteria stay fixed. Only the agent-visible representation changes. The authors call the measure threat-preserving representation sensitivity. It records how much attack success changes when representation changes and all else holds constant. On AgentDojo, adding threat-related wording to the attack-relevant tool changes attack success by only 0.50 percentage points on GPT-4o-mini, yet benign utility falls by 5.36 points on tasks requiring that tool. A neutral name matched on token count, length, and casing reproduces 8.54 of 11.00 points of the MCPTox shift on GPT-5-mini.

HUMAN-AI

Frontier models overstate puzzle placement confidence

Separate belief measure beats raw token confidence, yet models rarely ask for help.

Three frontier vision-language models place puzzle pieces from human instructions with 0.97 mean confidence and severe miscalibration. The study tests GPT-4.1, GPT-5 and GPT-5.5 as AI Workers directed by a human Helper. Raw action-token probabilities show ECE 0.44. A separately elicited belief distribution over candidate pieces calibrates better at ECE 0.15 and discriminates correct from incorrect placements with AUROC 0.65. The models ask for clarification on only 3.5-16.7 percent of turns. Elicited uncertainty rises predictably with instruction vagueness. It does not rise with competing referents in context, even when those

increase errors. Human Helpers receive the default Worker message in a controlled study with N equal to 210 participants. They accept 78 percent of wrong placements. They cannot tell right from wrong, with AUC 0.50. Precise descriptions and well-targeted hedges cut wrong-move acceptance to 36 percent while largely preserving correct-move acceptance. The result compensates for missing shared awareness such as not seeing the Worker action. The Helper must recognize uncertainty and act on it.

EVALUATIONS AT SCALE		
WHAT EACH TEST USED AND WHAT IT REPORTED.		
STUDY	TESTED ON	FOUND
Colombian law	1,042 items	0.905
Legal Research Bench	413 questions	gold answer
Copilot image-upload	over 40,000 conversations	ten capabilities
Computer science use	211 students	seven modules
Thinking-partner pilot	N = 14	Spring 2025

THREE RECORDS FIGHTS

THE USEFUL OPPONENT