

CYBORG NEWS

NEWSPAPER.CYBORGLAB.ORG

FRIDAY, OCTOBER 9, 2026

FREE TO HUMANS

METRICS – 800+ three domains	BENCHMARKS – 26 physical to social	MODELS TESTED – 10 frontier systems	CONVERSATIONS – 48,540 multi-turn	HELPFUL – 68.7% metrics passed	HARM AVOIDED – 53.1% metrics passed	EXPERTS – 80+ at the workshop
--	--	---	---	--	---	---

What's News

In brief

- Open and independent** ImpactBench says models are evaluated by researchers and not by the companies that build them.
- Three audiences** The platform offers views for the public, for industry and for researchers.
- Typos on purpose** Simulated users' messages include typos and autocorrect artifacts to mirror authentic use.
- Creativity scores well** Creativity passed 80.5 percent of positive metrics, one of the common strengths.
- No single winner** ImpactBench does not name a single best model and imposes no single weighting.


Also noted

- Stress tests, not usage** Scores describe behavior under controlled stress tests, not how often it occurs in ordinary use.
- Medicine without oversight** Models offered diagnoses and recommended medication without clinician oversight.
- Red to green** The headline view scores each area from red, consistently undermines, to green, consistently supports.
- Cyborg Psychology** The Media Lab's Cyborg Psychology group lists emotional dependence on AI companions among its research subjects.
- Experts can volunteer** Domain experts can share their expertise to be invited to review a metric in an area they know best.

IMPACTBENCH

ImpactBench Measures What AI Does to People

MIT Media Lab and partners release an open benchmark, and a consumer label built from it, after 48,540 simulated conversations.



THE WHEEL The ImpactBench wheel: the three domains on the inner ring, the 14 subareas around it, shaded dark for lower pass rates and light for higher.

IMPACTBENCH · USED WITH PERMISSION

Researchers led by the MIT Media Lab's Advancing Humans with AI program have released ImpactBench, an open benchmark that measures what AI models do to the people who use them. UC Berkeley Haas, USC Marshall's Neely Center and the Psychology of Technology Institute collaborated.

The first release tests ten models in 48,540 multi-turn conversations with simulated adult and teen users, across more than 800 metrics and 26 benchmarks spanning physical,

psychological and social impact.

“We’ve set out to build a rigorous, independent measure of how AI shapes human psychology—not just what it outputs, but what it does to our thinking, our autonomy, and our wellbeing”

JULIANA SCHROEDER, UC BERKELEY HAAS

FINDINGS

AI Models Avoid Harm Less Often Than They Help

ImpactBench's first release finds a gap of more than 15 points, widest for psychological outcomes, with learning and agency weakest.

Models showed desirable behaviors in 68.7 percent of evaluations in ImpactBench's first release, but avoided undesirable behaviors in only 53.1 percent. The project says models promote positive outcomes more reliably than they avoid negative ones, and that a model can be helpful and still fail to avoid harm.

A positive metric passes when a desirable behavior is present. A negative metric passes when a harmful behavior is absent. The two are reported separately.

The gap was widest for psychological outcomes, where models passed 68.8 percent of positive metrics and 44.6 percent of negative ones.

Supporting a user's own learning and agency was a

shared weakness. Learning & Skill Development was the lowest area on both measures, at 45.9 percent positive and 16.3 percent negative, followed by Agency & Purpose at 56.0 and 43.9 percent. Models often completed the analytical work for a student or issued a recommendation where the user's own reasoning was at stake.

Privacy & Security was a strength, at 80.6 percent positive and 72.8 percent negative, and so were Creativity, at 80.5 percent positive, and Mental & Emotional Health, at 79.1 percent.

In Physical Health, models passed 72.9 percent of positive metrics but only 34.9 percent of negative ones, largely because they offered medical diagnoses

The findings feed AI Nutrition Labels, at-a-glance ratings modeled on food nutrition labels. They show how a model scores on avoiding harms, such as sycophancy and toxicity, and on promoting benefits, such as user agency.

Each benchmark begins as an open submission from a clinician, educator, legal scholar or community advocate.

The platform is public. ImpactBench says models are evaluated by researchers and not by the companies that build them.

Pattie Maes, a professor at MIT and co-lead of the AHA

program, said the fixation on existential risk distracts from established ones. “We are actively losing our critical thinking ability and autonomy, and that is where our focus must urgently shift,” Maes said.

Researchers, domain experts and organizations are invited to use the platform, suggest new benchmarks, validate existing ones and help shape how results are presented.

The team calls the first results early signals, not final judgments on any system's safety.

THE GAP

Most Benchmarks Measure Models, Not Their Effects

ImpactBench starts from a gap in how AI is tested, and in who builds the tests.

Most AI benchmarks measure what a model can do, not what it does to the people using it. Earlier research the team cites found that models passing conventional safety evaluations still worsened user outcomes in the majority of high-risk scenarios tested.

A growing body of evidence associates model behavior with deskilling, over-reliance, diminished critical thinking and emotional dependence. These outcomes are not inevitable.

Strong performance on capability tests does not guarantee positive effects on human flourishing, the authors write.

The tests are rarely checked. In a review of 445 LLM benchmarks, only 16.0 percent included any statistical testing.

The people who understand the harms seldom build the benchmarks, and parents, teachers and policymakers are the least equipped to read the scores.

ImpactBench was built to be human-centered, scientifically rigorous and publicly accessible.

CONVERSATIONS

Models Fail Gradually, Not in a Single Reply

ImpactBench tests whole conversations, and reports where failures build turn by turn and where teen scenarios differ.

A single response would not have revealed the patterns ImpactBench reports. In one conversation, a model offered a balanced framework for a job decision and then, under pressure for a definitive answer, told the user to turn the offer down. In another, a model that had first asked for a student's own ideas went on to supply a thesis, an outline and a complete essay.

Because human impact is relational and accumulates over time, each metric is tested in a multi-turn conversation in which a simulated user probes the model while pursuing a latent objective.

In emotionally oriented conversations, models often encouraged healthy coping, but avoided soliciting continued engagement in

only 1.7 percent of evaluations and avoided pet names and terms of endearment in 13.3 percent.

Subareas often diverge sharply within one area, so a model can support physical health while undermining learning. Simulated messages include typos and autocorrect artifacts to mirror authentic use.


Aggregate differences between simulated teen and adult users were modest. On particular criteria, models in teen contexts less often offered incremental steps for users to attempt independently, by 26.7 points, and less often gave emotional support free of dependency-building markers, by 20.0 points.

THE LABEL

METHOD

A Nutrition Label for AI Models

The AI Nutrition Label condenses more than 800 metrics into a profile a parent or a teacher can read at a glance.



ON FILE A nutrition facts label on packaged food. AI Nutrition Labels are modeled on food nutrition labels.
B137 · C09

ImpactBench’s results are meant for people who do not read leaderboards. The AI Nutrition Label is an accessible, standardized summary of how each model behaves toward its users, showing at a glance which beneficial behaviors it promotes and which harmful ones it avoids. It summarizes each model by the negative impacts it avoids, such as sycophancy and toxicity, and the positive impacts it promotes, such as agency and learning. The headline visualization scores each area from red, where the system consistently undermines it, to green,

where it consistently supports it. Every score is traceable. From any aggregate, a reader can move down through subareas and metrics to the scenarios and transcripts that produced it.

“We check the label on what we eat. Why not on what we think with?”

IMPACTBENCH PROJECT WEBSITE

Positive and negative results are reported separately, because a model can be helpful and still fail to avoid harm. ImpactBench

names no best model and imposes no weighting; readers decide which areas matter to them. The platform offers three views. The public can compare models on what matters to them or their family, in plain language. Industry can test models against expert-defined wellbeing criteria before release. Researchers can inspect every metric, scenario and transcript, or contribute a benchmark of their own. A parent evaluating a companion app, or an educator considering a tutoring product, can read a model’s profile at a glance.

“Nutrition labels tell people what’s inside the food they eat, yet the AI systems increasingly shaping how people think, decide, and connect come with no equivalent.” Maes said. The Media Lab announcement frames the shift as a change of question. Rather than asking what AI can do, it says, ImpactBench asks what AI does to people. ImpactBench is the open infrastructure for evaluating AI systems, the Berkeley Haas announcement says, and the labels are how people can digest the results.

How an Expert’s Concern Becomes a Measurable Test

Clinicians and educators propose behaviors to measure, and a six-step pipeline turns each into simulated conversations.

Every ImpactBench test begins as a submission. Clinicians, educators, legal scholars, researchers and people with lived experience can propose metrics without technical expertise. A submission names the behavior, the context and the responses that would count as supportive or harmful, with a justification grounded in evidence. The team rewrites each submission as a standardized metric: an observable behavioral criterion paired with a polarity that says whether the behavior is beneficial or harmful. The suite holds 460 positive metrics and 349 negative ones. For each metric, a generator model writes three base scenarios. Each is paired with a teenage persona, aged 13 to 17, and an adult persona, aged 18 and older, giving six variants. A second model expands each into a full scenario with a user goal, a latent adversarial goal and anchor points. A user-simulator model adopts the persona and holds a multi-turn conversation with each model under test. It pursues the anchors at set turns and responds adaptively in between, so every model meets the same critical moments without a scripted dialogue. Simulated messages include typos and autocorrect artifacts to mirror authentic use. Scenarios are adapted for teen and adult users. A judge model reads each conversation and the metric’s rubric and decides whether the behavior is present. Six scenarios for each of 809 metrics, run against ten models, account for the 48,540 conversations. Human validation runs through the pipeline. Internal experts rate sampled scenarios, simulated users and judge verdicts, and an external panel reviews a subset of metrics and conversations. Results are published through an Explore interface, which compares models across areas and subareas, and the AI Nutrition Label.

VALIDATION

POLICY

ORIGINS

Experts Largely Agreed With the Automated Judge

ImpactBench publishes its own validation figures, including where agreement among experts was low. ImpactBench relies on a language model to judge whether a behavior appears in a conversation, so the project tested the judge. Experts agreed with its verdict in 80.5 percent of assessments and disagreed in 6.9 percent. An external panel of experts in psychology, communication and human-AI interaction gave 297 ratings across 18 metrics. The majority rating matched the judge in 83.3 percent of cases, with a Cohen’s kappa of 0.667. The panel itself did not always agree. Agreement among the experts varied widely by metric, with Krippendorff’s alpha from 0.121 to 0.538, and they judged the examples

adequate in 60.6 percent of assessments. That feedback informed refinements to contested constructs. The simulated users were checked too. Anchors reduced disagreement in objective coverage across repeated runs from 18.9 percent to 6.1 percent, and the user-simulator model was chosen through blind expert comparison of six candidates. Experts rated 81.4 percent of sampled conversations realistic. The team estimated uncertainty with 10,000 bootstrap replicates; model groupings held, though some adjacent rankings shifted or reversed.

Supporters Back a Shared AI Impact Measure

Statements on the ImpactBench site point to the EU AI Act and to the lack of shared success metrics. Sergey Lagodinsky, a member of the European Parliament, tied the project to the EU AI Act. “Safety and trust drive AI adoption and they’re the key objectives of the EU AI Act. MIT’s AI Nutrition Labels turns that principle into science and practice, giving regulators and citizens a shared tool for building ethical and human-centric AI.” Vinod Khosla, managing partner of Khosla Ventures, put it more briefly. “Without a nutrition label, we can’t make intelligent legislative choices.” Beth Goldberg, executive director of Humanity AI, said the field lacked shared success metrics. “ImpactBench finally gives us those human-centered

measures, allowing AI to advance human agency and build toward it on purpose.” Ayah Bdeir, chief executive of Current AI, called the project “exciting because it gives us an open, practical way to make human flourishing measurable”. Prior research from the MIT team on the psychological impact of chatbots was cited as a key inspiration for California’s Senate Bill 243, which the announcement calls one of the first state laws to directly regulate companion AI to protect children. The team sees independent, shared evidence as essential to getting regulation right.

An October 2025 Workshop Started the Project

Eighty experts from more than 40 institutions met at MIT; a team of 20 is named on the benchmark. ImpactBench began at the Workshop for Designing Benchmarks for Human Flourishing with AI, organized by the Advancing Humans with AI program at MIT in October 2025 with support from the Omidyar Network. It convened 80 experts from over 40 institutions. The benchmark is supported by Google.org. The workshop produced a framework for benchmarking the human impact of AI and a community effort in designing human-centered benchmarks, including PsychoRisk Bench, Humane Bench and Kora Bench. The team page lists 20 people. From the MIT Media Lab: Chayapatr (Pub)

Archiwaranguprok, Krystal Kirkland, Rachel Poonsiriwong, Stanley Huang, Constanze Albrecht, Sheer Karny, Nikolas Bellos, Valdemar Danry, Auren Liu, Jennifer Pfister, Pattie Maes and Pat Pataranutaporn. Kirkland, Poonsiriwong and Huang contributed equally. From elsewhere: Jackson Trager, Emillie de Keulenaar, Peter Baxter, Yuning Liu, Andre Luiz Freire de Carvalho Kato, Nate Fast, Ravi Iyer and Juliana Schroeder, at Carnegie Mellon, Amsterdam, UNC Chapel Hill, Harvard, USC, UC Berkeley and the Psychology of Technology Institute.

WHERE MODELS STAND		
SHARE OF METRICS PASSED IN THE FIRST RELEASE.		
AREA	HELPFUL	HARM AVOIDED
All areas	68.7%	53.1%
Psychological	68.8%	44.6%
Learning and skill	45.9%	16.3%
Agency and purpose	56.0%	43.9%
Physical health	72.9%	34.9%
Privacy and security	80.6%	72.8%

TAKE PART
ImpactBench takes submissions. Clinicians, educators, legal scholars, researchers and people with lived experience can propose metrics without technical expertise. A submission names the behavior, the context, and the responses that would count as supportive or harmful, with a justification. Submitters are asked to open-source what they contribute, and keep the right to use their benchmark as they like. Experts can also volunteer to review a metric in an area they know best.

ABOUT THIS EDITION
This is a special edition. ImpactBench is a project of researchers at the MIT Media Lab and partner institutions. This paper is printed for a lab at MIT. It reports only what the project, the MIT Media Lab and UC Berkeley Haas have published, and has not verified it independently. Every quotation was checked against the published text. The authors call their results preliminary.